

Daily Learning Report: Understanding LSTMs for Time Series Prediction

Date: [Insert Date]
Author: [Your Name]


1. Introduction

Today’s focus was on understanding **Long Short-Term Memory (LSTM)** networks, a type of Recurrent Neural Network (RNN) designed to handle sequential data like stock prices, weather data, or text. Below is a summary of key concepts, diagrams, and takeaways.

2. What is an LSTM?

An LSTM is a neural network with **memory cells** that can retain information over long periods. It uses **gates** to control what information to keep, forget, or pass to the next timestep.

LSTM Cell Structure:



Key:
 σ : Sigmoid (squishes values to 0–1).
 \tanh : Hyperbolic tangent (squishes values to -1–1).

3. How LSTMs Process Sequences

For time series prediction, LSTMs process sequences using **sliding windows**.

Example: Predicting Day 4 from Days 1–3




Hidden State (h_t): Passed to the next timestep (short-term memory).

Cell State (C_t): Carries long-term dependencies (e.g., trends).

4. Key Concept: Parameter Sharing

All LSTM cells in a layer **share the same weights and biases** (W_f , W_i , W_C , W_o , b_f , etc.).

Diagram: Shared Weights Across Windows




Updated Weights → Updated Weights → Final Weights (W_n , b_n)

Final Weights (W_n , b_n): Used to predict test data (e.g., Days 8–10 → Day11).

5. Training Process


Step 1: Forward Pass
Process the sequence (e.g., Days 1–3) to predict the next value (Day4).
Compute loss (e.g., Mean Squared Error).
Step 2: Backward Pass (BPTT)
Calculate gradients for shared weights using **Backpropagation Through Time (BPTT)**.
Update weights to minimize loss.
Step 3: Repeat for All Windows
Train on all windows to refine weights iteratively.

Diagram: Gradient Accumulation



6. Testing the Model
Use the **final weights (W_n, b_n)** from training to predict unseen sequences.
Example: Predict Day11 using Days 8–10.

Test Data Flow



7. Common Misconceptions

Myth	Fact
Each window has separate weights.	All windows share the same weights (W, b).
Training on one window is enough.	Train on all windows to generalize.
Hidden states carry across windows.	Hidden states reset for each new window.

8. Key Takeaways

- . **Parameter Sharing:** LSTMs reuse weights across all timesteps/windows.
- . **Gates Control Memory:** Forget/Input/Output gates manage information flow.
- . **Training:** Updates weights iteratively using gradients from all windows.
- . **Prediction:** Final weights (W_n, b_n) predict unseen sequences.

9. Next Steps

Code Implementation: Try building an LSTM with TensorFlow/PyTorch.
Hyperparameter Tuning: Experiment with `look_back` window size, hidden units.
Advanced Topics: Bidirectional LSTMs, Attention Mechanisms.

End of Report

Let me know if you’d like a code example or further details! 😊